

CME 213

Lecture 21: Atomics & Segmented Scan

David Tarjan, NVIDIA Research

Overview

- **Atomics**
- **Segmented Scan**

ATOMICS

The Problem

- How do you do global communication?
- Finish a grid and start a new one
- Scan
- **Atomics**

Race Conditions

- **Coordinate by writing to a predefined memory location**
 - **Race condition! Updates can be lost**

Race Conditions

threadId:0	threadId:1917
<code>// vector[0] was equal to 0</code>	
<code>vector[0] += 5;</code>	<code>vector[0] += 1;</code>
<code>...</code>	<code>...</code>
<code>a = vector[0];</code>	<code>a = vector[0];</code>

- What is the value of a in thread 0?
- What is the value of a in thread 1917?

Race Conditions

- Thread 0 could have finished execution before 1917 started
- Or the other way around
- Or both are executing at the same time

Race Conditions

- **Answer: not defined by the programming model, can be arbitrary**

Atomics

- CUDA provides **atomic** operations to deal with this problem

Atomics

- An atomic operation guarantees that only a single thread has access to a piece of memory while an operation completes
- The name atomic comes from the fact that it is uninterruptable
- No dropped data, but ordering is still arbitrary
- Different types of atomic instructions
- `atomic{Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor}`
- More types in fermi

Example: Histogram

```
// Determine frequency of colors in a picture
// colors have already been converted into ints
// Each thread looks at one pixel and increments
// a counter atomically
```

```
__global__ void histogram(int* color,
                          int* buckets)
{
    int i = threadIdx.x
          + blockDim.x * blockIdx.x;
    int c = colors[i];
    atomicAdd(&buckets[c], 1);
}
```

Example: Workqueue

```
// For algorithms where the amount of work per item  
// is highly non-uniform, it often makes sense for  
// to continuously grab work from a queue
```

```
__global__
```

```
void workq(int* work_q, int* q_counter,  
           int* output, int queue_max)  
{  
    int i = threadIdx.x  
           + blockDim.x * blockIdx.x;  
    int q_index =  
        atomicInc(q_counter, queue_max);  
    int result = do_work(work_q[q_index]);  
    output[i] = result;  
}
```

Atomics

- **Atomics are slower than normal load/store**
- **You can have the whole machine queuing on a single location in memory**
- **Atomics unavailable on G80!**

Example: Global Min/Max (Naive)

```
// If you require the maximum across all threads  
// in a grid, you could do it with a single global  
// maximum value, but it will be VERY slow
```

__global__

```
void global_max(int* values, int* gl_max)  
{  
    int i = threadIdx.x  
        + blockDim.x * blockIdx.x;  
    int val = values[i];  
    atomicMax(gl_max, val);  
}
```

Example: Global Min/Max (Better)

```
// introduce intermediate maximum results, so that
// most threads do not try to update the global max
__global__
void global_max(int* values, int* gl_max,
               int *reg_max,
               int num_regions)
{
    // i and val as before ...
    int region = i % num_regions;
    if(atomicMax(&reg_max[region], val) < val)
    {
        atomicMax(gl_max, val) ;
    }
}
```

Global Min/Max

- **Single value causes serial bottleneck**
- **Create hierarchy of values for more parallelism**
- **Performance will still be slow, so use judiciously**
- **See next lecture for even better version!**

Summary

- Can't use normal load/store for inter-thread communication because of **race conditions**
- Use **atomic instructions** for sparse and/or unpredictable global communication
 - Scan is good for dense communication pattern and where ordering is needed
- **Decompose data** (very limited use of single global sum/max/min/etc.) for more parallelism

SEGMENTED SCAN

Segmented Scan

- **What it is:**
 - **Scan + Barriers/Flags associated with certain positions in the input arrays**
 - **Operations don't propagate beyond barriers**
- **Do many scans at once, no matter their size**

Head flag



Index if head element 0 otherwise



Max scan

mindex



index



Data



Data



Data



Data



Image taken from
“Efficient parallel
scan algorithms
for GPUs” by S.
Sengupta, M.
Harris, and M.
Garland

Segmented Scan

```
__global__ void segscan(int * data,  
    int * flags)  
{  
    __shared__ int s_data[BL_SIZE];  
    __shared__ int s_flags[BL_SIZE];  
    int idx = threadIdx.x + blockDim.x  
    * blockIdx.x;  
    // copy block of data into shared  
    // memory  
    s_data[idx] = ...; s_flags[idx] = ...;  
    __syncthreads();
```

Segmented Scan

...

```
// choose whether to propagate
```

```
s_data[idx] = s_flags[idx] ?
```

```
    s_data[idx] :
```

```
    s_data[idx - 1] + s_data[idx];
```

```
// create merged flag
```

```
s_flags[idx] =
```

```
    s_flags[idx - 1] | s_flags[idx];
```

```
// repeat for different strides
```

```
}
```

Segmented Scan

- **Doing lots of reductions of unpredictable size at the same time is the most common use**
- **Think of doing sums/max/count/any over arbitrary sub-domains of your data**

Segmented Scan

- **Common Usage Scenarios:**
 - Determine which region/tree/group/object class an element belongs to and assign that as its new ID
 - Sort based on that ID
 - Operate on all of the regions/trees/groups/objects in parallel, no matter what their size or number

Segmented Scan

- Also useful for implementing divide-and-conquer type algorithms
 - Quicksort and similar algorithms

Questions?

Backup Slides

Example Segmented Scan

```
int data[10] = {1, 1, 1, 1, 1, 1, 1, 1, 1, 1};
```

```
int flags[10] = {0, 0, 0, 1, 0, 1, 1, 0, 0, 0};
```

```
int step1[10] = {1, 2, 1, 1, 1, 1, 1, 2, 1, 2};
```

```
int flg1[10] = {0, 0, 0, 1, 0, 1, 1, 1, 0, 0};
```

```
int step2[10] = {1, 2, 1, 1, 1, 1, 1, 2, 1, 2};
```

```
int flg2[10] = {0, 0, 0, 1, 0, 1, 1, 1, 0, 0};
```

```
...
```

Example Segmented Scan

```
int step2[10] = {1, 2, 1, 1, 1, 1, 1, 2, 1, 2};
```

```
int flg2[10]  = {0, 0, 0, 1, 0, 1, 1, 1, 0, 0};
```

```
...
```

```
int result[10] = {1, 2, 3, 1, 2, 1, 1, 2, 3, 4};
```