

# Sparse Matrix-Vector Multiplication & Texture

## Lesson 22

David Tarjan (most slides from  
Nathan Bell's presentation on  
SpMV)



# Overview




- **Sparse Matrix Vector Product**
  - Performance Considerations
  - Matrix Formats
- **Texture**
  - Overview
  - Example Usage for SpMV

# Dense Matrix-Vector Multiplication



- **SGEMV / DGEMV in BLAS**
  - Memory-bound performance



Dense Matrix

# Sparse Matrix-Vector Multiplication



- One multiply-add per nonzero entry
  - Some reuse of vector data



# Performance



- Performance is *Memory-Bound*
  - 5-20 GFLOP/s is typical
- Low Arithmetic Intensity
  - 2 flops : 8+ bytes (float)
  - 2 flops : 16+ bytes (double)
- Primary objective
  - Use memory bandwidth efficiently

# Performance



- **Tesla C2050 floating point performance**
  - Single 1,030 GFLOP/s (peak)
  - Double 515 GFLOP/s (peak)
- **Tesla C2050 memory bandwidth**
  - 144 GB/s (peak)
- **Intensity Threshold**
  - 7.14 FLOP : Byte (single)
  - 3.57 FLOP : Byte (double)

- **Tesla C2050 threshold**
  - 7.14 FLOP : Byte (single)
  - 3.57 FLOP : Byte (double)
- ***Dense* Matrix-Vector Multiplication Intensity**
  - 0.25 - 0.50 FLOP : Byte (single)
  - 0.12 - 0.25 FLOP : Byte (double)
- ***Sparse* Matrix-Vector Multiplication Intensity**
  - 0.12 - 0.50 FLOP : Byte (single)
  - 0.08 - 0.25 FLOP : Byte (double)

# Performance Considerations




- **Use memory bandwidth efficiently**
  - Memory coalescing
- **Expose lots of fine-grained parallelism**
  - Need thousands of threads
- **Find opportunities for reuse**
  - Make use of caching



# Memory Coalescing




- Recall: *blocks* divided into physical *warps*



# Memory Coalescing




- Fully Coalesced Memory Access



# Memory Coalescing



- **Partially Coalesced Memory Access**



# Memory Coalescing




- Misaligned memory access



# Memory Coalescing



- **Uncoalesced Memory Access**
  - Separated by 32+ words



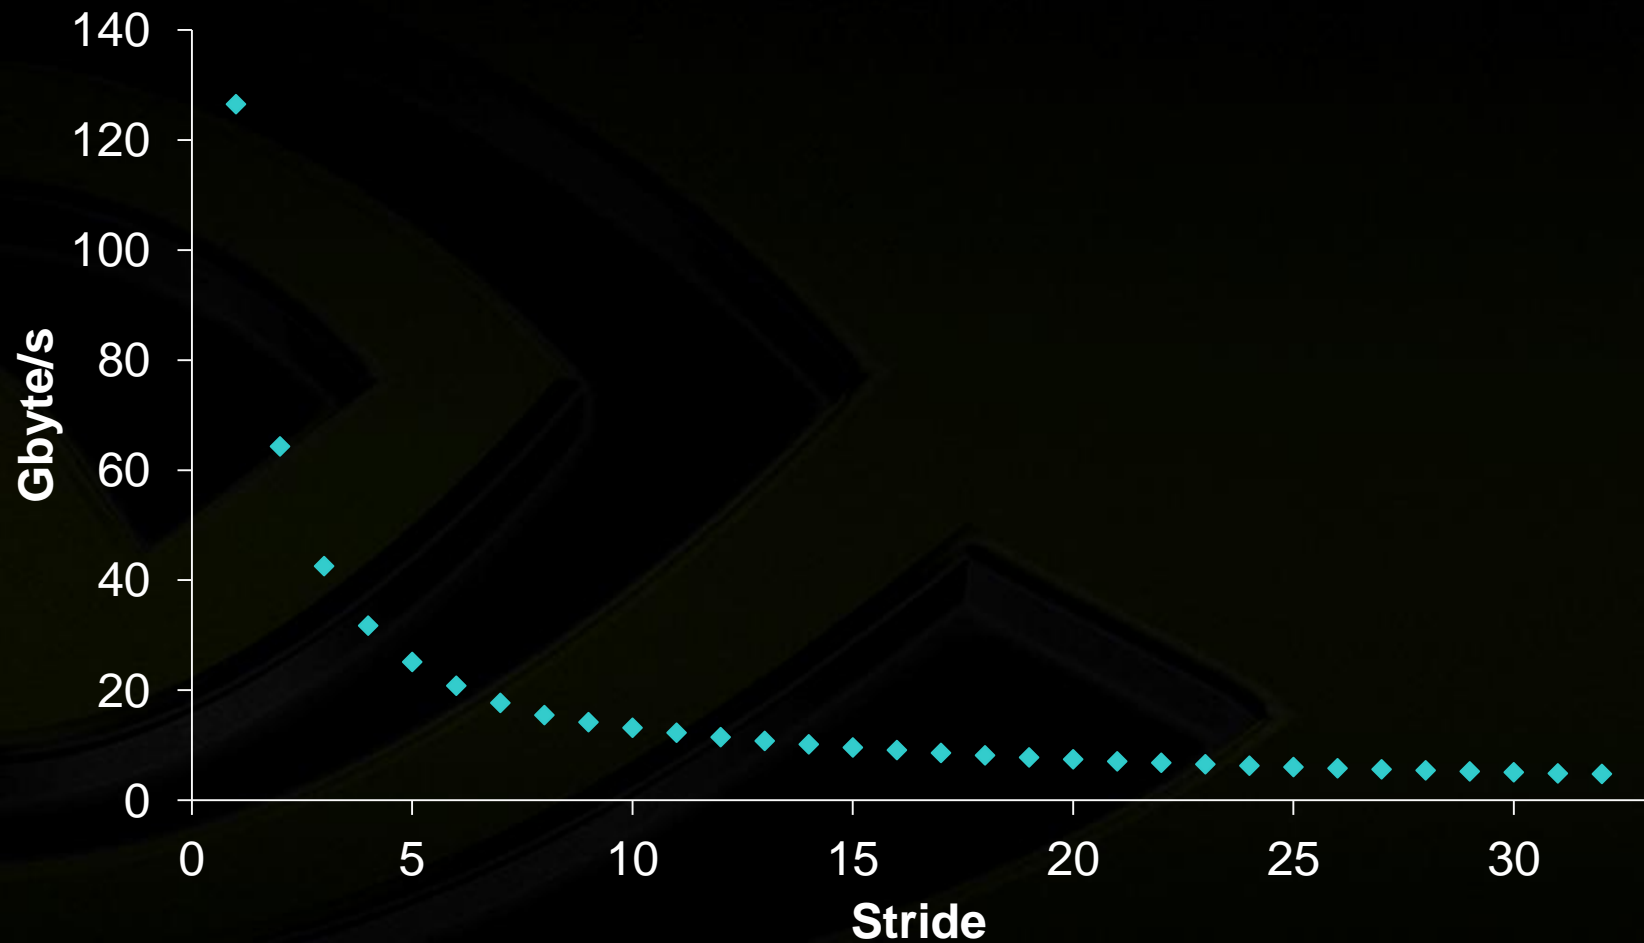
# Memory Coalescing (SAXPY)



- Example: SAXPY with stride
  - Fully Coalesced when stride is 1

```
for (int i = 0; i < N; i++)  
    y[stride * i] += a * x[stride * i];
```

# Memory Coalescing (SAXPY)



# Memory Coalescing (SAXPY)



- **Example: SAXPY with offset**
  - Aligned when `offset` is 0

```
for (int i = 0; i < N; i++)  
    y[i + offset] += a * x[i + offset];
```



# Memory Alignment (SAXPY)



# Types of Memory Access

- **Reading matrix structure**
  - Determined by matrix format
  - Most bandwidth consumption
- **Reading source vector (x)**
  - Determined by matrix structure
  - Little control over access pattern
  - Potential reuse
- **Writing destination vector (y)**
  - Little bandwidth consumption

# Compressed Sparse Row (CSR)

- Rows laid out in sequence
- Inconvenient for fine-grained parallelism



# CSR SpMV (serial)



```
void csr_spmv(int      num_rows,
              int      * row_offsets,
              int      * column_indices,
              float * values,
              float * x,
              float * y)
{
    for (int row = 0; row < num_rows; row++)
    {
        int row_begin = row_offsets[row];
        int row_end   = row_offsets[row + 1];

        float sum = 0;

        for (int offset = row_begin; offset < row_end; offset++)
            sum += values[offset] * x[column_indices[offset]];

        y[row] = sum;
    }
}
```

# CSR (scalar) kernel



```
__global__
void csr_spmv(int      num_rows,
              int      * row_offsets,
              int      * column_indices,
              float * values,
              float * x,
              float * y)
{
    int row = blockDim.x * blockIdx.x + threadIdx.x;

    if (row < num_rows)
    {
        int row_begin = row_offsets[row];
        int row_end   = row_offsets[row + 1];

        float sum = 0;


        for (int offset = row_begin; offset < row_end; offset++)
            sum += values[offset] * x[column_indices[offset]];

        y[row] = sum;
    }
}
```

# CSR (scalar) kernel




- One thread per row
  - Poor memory coalescing
  - Unaligned memory access



# CSR (vector) kernel




- One SIMD vector or *warp* per row
  - Partial memory coalescing
  - Unaligned memory access



# CSR (vector) kernel

- Reduce partial sums in shared memory
  - Example: warp of 8 threads



Shared memory



# CSR (vector) kernel



```
__global__ void
spmv_csr_vector_kernel()
{
    __shared__ volatile ValueType sdata[VECTORS_PER_BLOCK * THREADS_PER_VECTOR];
    __shared__ volatile IndexType ptrs[VECTORS_PER_BLOCK][2];

    // use two threads to fetch Ap[row] and Ap[row+1]
    if(thread_lane < 2) { ptrs[vector_lane][thread_lane] = Ap[row + thread_lane]; }

    // implicit synchronization here, no synchronization due to warp synchronous behavior assumed
    const IndexType row_start = ptrs[vector_lane][0];           //same as: row_start = Ap[row];
    const IndexType row_end   = ptrs[vector_lane][1];           //same as: row_end   = Ap[row+1];

    // initialize local sum
    ValueType sum = 0;

    // accumulate local sums
    for(IndexType jj = row_start + thread_lane; jj < row_end; jj += THREADS_PER_VECTOR)
        sum += Ax[jj] * x[Aj[jj]];

    // store local sum in shared memory
    sdata[threadldx.x] = sum;



    // reduce local sums to row sum
    <normal warpscan here>

    // first thread writes the result
    if (thread_lane == 0)
        y[row] = sdata[threadldx.x];
}
```

# ELL kernel



- Full coalescing




column-major ordering

# ELL kernel




- Pad columns for alignment



# DIA kernel




- Same as ELL



# COO kernel




- One thread per nonzero element
  - Fully coalesced



# COO kernel



- Store  $i$  and  $A(i, j) * x(j)$  in shared memory
  - Compute row sums using *segmented* reduction




# Memory Coalescing Summary



- **Full Coalescing**
  - DIA, ELL, and COO
- **Partial Coalescing**
  - CSR (vector)
  - Efficiency depends on row length
- **Little Coalescing**
  - CSR (scalar)

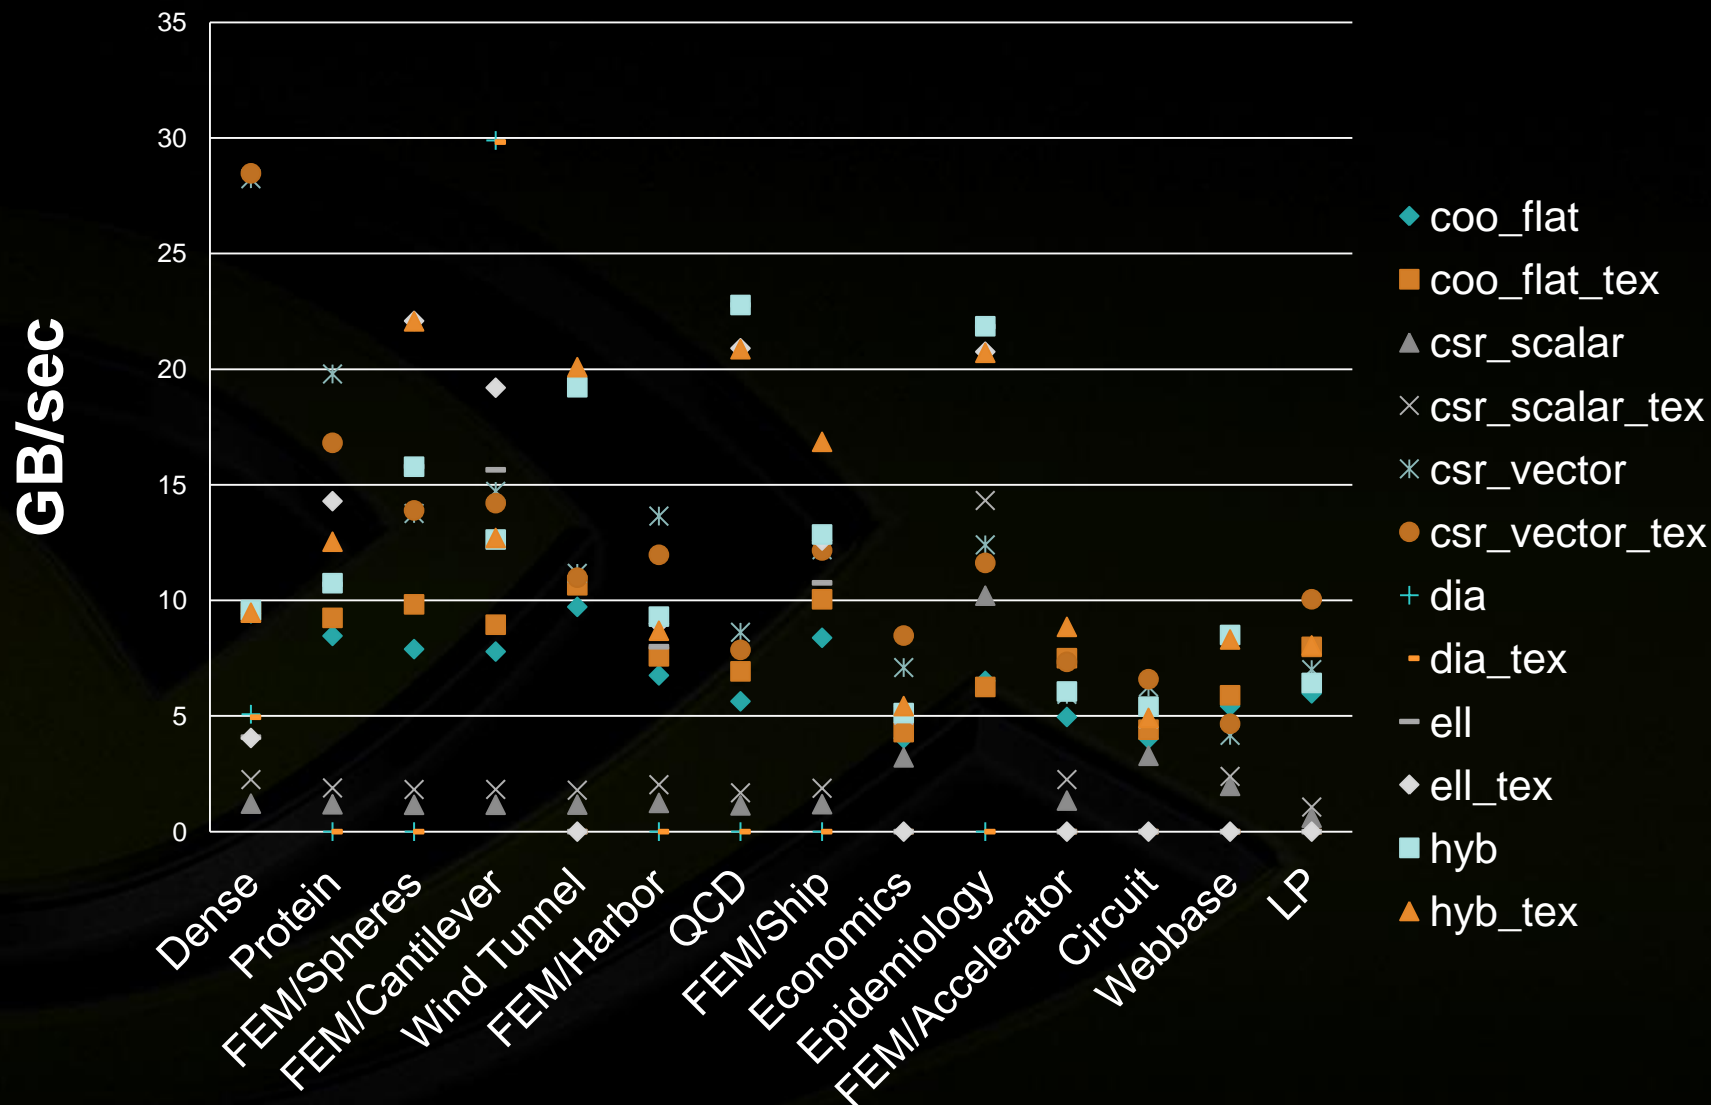
# Performance Comparison of Formats



Data collected on a C2075 by Steve Dalton with latest CUSP



# Performance Comparison of Formats



Data collected on a C2075 by Steve Dalton with latest CUSP


# Caching

- Repeated accesses to source vector



# Caching

- Effectiveness depends on matrix structure



# Caching




- **Fermi architecture has L1 cache**
  - No effort needed
- **Earlier architectures have texture cache**
  - Often worth ~30% improvement
- **Software-managed cache**
  - Preload into shared memory
- **Effectiveness depends on matrix structure**

# Other Techniques



- **Block Formats**

- Reduce index overhead



- **Multiple Vectors**

- Reuse matrix data



# Performance Considerations



- **Use memory bandwidth efficiently**
  - Memory coalescing
- **Expose lots of fine-grained parallelism**
  - Need thousands of threads
- **Find opportunities for reuse**
  - Make use of caching



# Generic Parallel Algorithms for Sparse Matrix and Graph Computations

## CUSP

# Example Usage



```
#include <cusp/coo_matrix.h>
#include <cusp/array1d.h>
#include <cusp/multiply.h>

int main(void)
{
    size_t M    = 10;
    size_t N    = 15;
    size_t NNZ  = 43;

    // allocate 10x15 COO matrix and vectors
    cusp::coo_matrix<int, float, cusp::device_memory> A(M, N, NNZ);
    cusp::array1d<float, cusp::device_memory> x(N);
    cusp::array1d<float, cusp::device_memory> y(M);

    // initialize A and x
    ...

    // compute matrix-vector product  $y = A * x$ 
    cusp::multiply(A, x, y);

    return 0;
}
```



# Algorithms



- **Multiply**
  - Sparse Matrix \* Vector
  - Sparse Matrix \* Sparse Matrix
- **Level 1 BLAS**
- **Transpose**
- **Maximal Independent Sets**
- **More to come**

# Sparse Matrix Containers



```
#include <cusparse/csr_matrix.h>

int main(void)
{
    // allocate storage for (4,3) matrix with 6 nonzeros
    cusparse::csr_matrix<int, float, cusparse::host_memory> A(4,3,6);

    // initialize matrix entries on host
    A.row_indices[0] = 0; A.column_indices[0] = 0; A.values[0] = 10.0f;
    A.row_indices[1] = 0; A.column_indices[1] = 2; A.values[1] = 20.0f;
    A.row_indices[2] = 2; A.column_indices[2] = 2; A.values[2] = 30.0f;
    A.row_indices[3] = 3; A.column_indices[3] = 0; A.values[3] = 40.0f;
    A.row_indices[4] = 3; A.column_indices[4] = 1; A.values[4] = 50.0f;
    A.row_indices[5] = 3; A.column_indices[5] = 2; A.values[5] = 60.0f;

    // A now represents the following matrix
    //      [10  0 20]
    //      [ 0  0  0]
    //      [ 0  0 30]
    //      [40 50 60]

    return 0;
}
```

# Sparse Matrix Containers



- **COO – Coordinate format**
- **CSR – Compressed Sparse Row Format**
- **DIA – Diagonal Format**
- **ELL – ELLPACK Format**
- **HYB – Hybrid ELL + COO Format**

# TEXTURE

# Texture Use



```
texture<float, 1, cudaReadModeElementType> t_foo;

__global__ bar_kernel(float * d_bar)
{
    int index = ...
    float fromTex = tex1D(t_foo,
index);
    float fromArray = d_bar[index];
}
```

# Texture Binding



```
//create cuda array
    cudaChannelFormatDesc channelDesc =
cudaCreateChannelDesc(32, 0, 0, 0, cudaChannelFormatKindFloat);
    cudaMallocArray( &d_foo_array, &channelDesc, size, 1 );
    cudaMemcpyToArray( d_foo_array, 0, 0, d_foo, size * sizeof
(float), cudaMemcpyDeviceToDevice);

// Bind the array to the texture
cudaBindTextureToArray( t_foo, d_some_array, channelDesc);
```

# CSR (vector) kernel with texture



```
texture<float, 1, cudaReadModeElementType> t_x;
```

```
__global__ void
spmv_csr_vector_tex_kernel()
{
    ...
    // initialize local sum
    ValueType sum = 0;

    // accumulate local sums
    for(IndexType jj = row_start + thread_lane; jj < row_end; jj +=
THREADS_PER_VECTOR)
        sum += Ax[jj] * tex1D(t_x, Aj[jj]); // access to x is a sparse gather

    // store local sum in shared memory
    sdata[threadIdx.x] = sum;

    ...
}
```

# References



## *Implementing Sparse Matrix-Vector Multiplication on Throughput-Oriented Processors*

Nathan Bell and Michael Garland

Proceedings of Supercomputing '09

## *Efficient Sparse Matrix-Vector Multiplication on CUDA*

Nathan Bell and Michael Garland

NVIDIA Technical Report NVR-2008-004, December 2008

## *Model-driven Autotuning of Sparse Matrix-Vector Multiply on GPUs*

Jee Whan Choi, Amik Singh and Richard W. Vuduc

Proceedings of Principles and Practice of Parallel Programming (PPoPP) 2010



# Backup Slides



# Performance Considerations



- **Use memory bandwidth efficiently**
  - Memory coalescing
- **Expose lots of fine-grained parallelism**
  - Need thousands of threads
- **Find opportunities for reuse**
  - Make use of caching



# Exposing Parallelism



- **DIA, ELL & CSR (scalar)**
  - One thread per row
- **CSR (vector)**
  - One warp per row
- **COO**
  - One thread per nonzero



Finer Granularity


# Performance Considerations



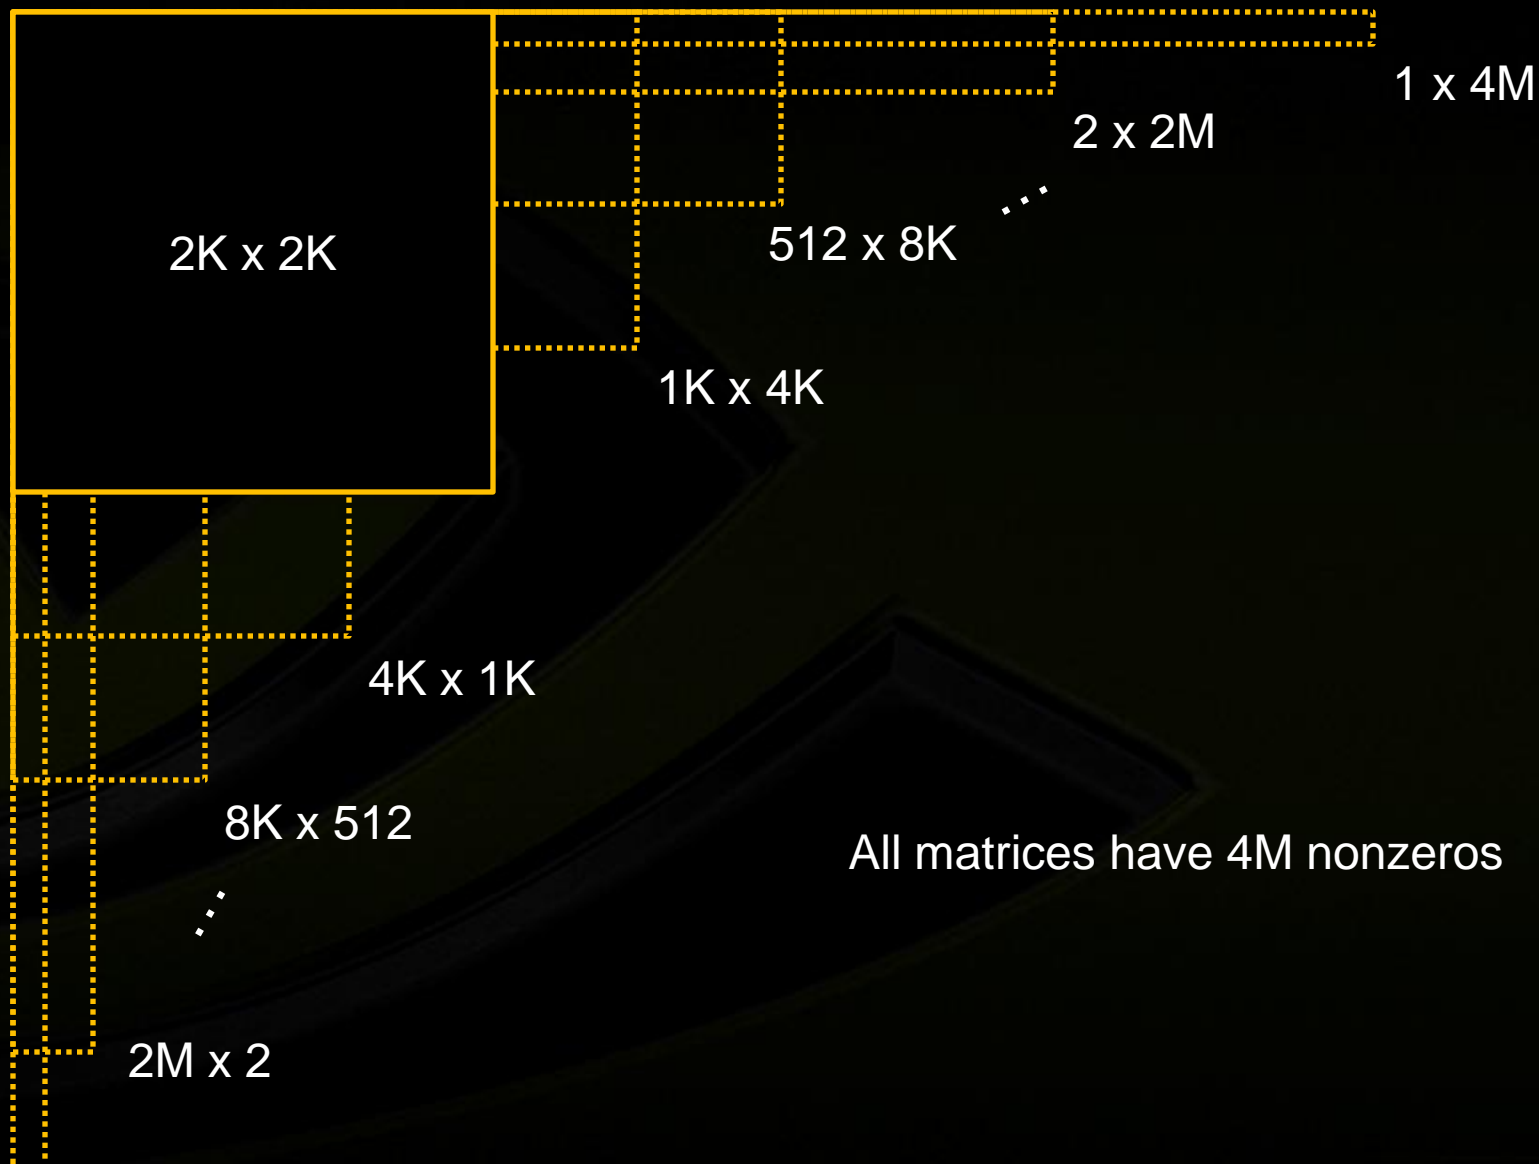
- **Use memory bandwidth efficiently**
  - Memory coalescing
- **Expose lots of fine-grained parallelism**
  - Need thousands of threads
- **Find opportunities for reuse**
  - Make use of caching




# Performance Comparison of Formats



# Exposing Parallelism



# Exposing Parallelism



# Exposing Parallelism



- **One thread per row**
  - ELL, DIA, and CSR (scalar) kernel
  - Generally good enough (>20K rows is common)
- **One warp per row**
  - CSR (vector) kernel
  - Fewer rows is sufficient (>256)
- **One thread per entry**
  - COO kernel
  - Insensitive to matrix shape