

# CME213/ME339

## Lecture 12

Eric Darve  
Erich Elsen  
Austin Gibbons

Department of Mechanical Engineering  
Institute for Computational and Mathematical Engineering  
Stanford University

Spring 2012



# Thrust Algorithms

---

- Transformations
- Reductions
- Sorts
- Searches
- Set Operations
- Compaction



# Algorithm variants

---

- Many algorithms will have many variants on the same idea
- Often there can be either 1 or 2 input sequences with either an unary or binary function
- \*\_if variants take a stencil or predicate function and only perform the operation if the stencil or predicate is true
- \*\_by\_key variants perform the operation on multiple smaller sequences
  - with the exception of sort\_by\_key which is poorly named



# Transform Algorithms

---

## transform

```
1 OutputIterator thrust::transform(InputIterator first,  
2                               InputIterator last,  
3                               OutputIterator result,  
4                               UnaryFunction op);
```

## Example use:

```
1 #include <thrust/transform.h>  
2 #include <thrust/functional.h>  
3  
4 int data[10] = {-5, 0, 2, -3, 2, 4, 0, -1, 2, 8};  
5 thrust::negate<int> op;  
6 thrust::transform(data, data + 10, data, op); // in-place transformation  
7 // data is now {5, 0, -2, 3, -2, -4, 0, 1, -2, -8};
```



# Transform Algorithms

## transform

```
1  OutputIterator thrust::transform(InputIterator1 first1,
2                                  InputIterator1 last1,
3                                  InputIterator2 first2,
4                                  OutputIterator result,
5                                  BinaryFunction op);
```

## Example use:

```
1  #include <thrust/transform.h>
2  #include <thrust/functional.h>
3
4  thrust::host_vector<int> input1(6);
5  // = {-5, 0, 2, 3, 2, 4}
6  thrust::host_vector<int> input2(6);
7  // = { 3, 6, -2, 1, 2, 3};
8  thrust::host_vector<int> output(6);
9
10 thrust::plus<int> op;
11 thrust::transform(input1.begin(), input1.end(), input2.begin(),
12                  output.begin(), op);
13 // output is now {-2, 6, 0, 4, 4, 7};
```



# Transform Algorithms

## transform\_if

```
1 ForwardIterator thrust::transform_if (InputIterator first,
2                                     InputIterator last,   ForwardIterator result,
3                                     UnaryFunction op,     Predicate pred);
```

## Example use:

```
1  #include <thrust/transform.h>
2  #include <thrust/functional.h>
3  int data[10]    = {-5, 0, 2, -3, 2, 4, 0, -1, 2, 8};
4
5  struct is_odd {
6      __host__ __device__
7      bool operator()(int x) {
8          return x % 2;
9      }
10 };
11
12 thrust::negate<int> op;
13 // negate odd elements; in-place transformation
14 thrust::transform_if(data, data + 10, data, op, is_odd());
15
16 // data is now {5, 0, 2, 3, 2, 4, 0, 1, 2, 8};
```



# Transform Algorithms

## transform\_if

```
1 ForwardIterator thrust::transform_if(InputIterator1 first1,
2   InputIterator1 last1,           InputIterator2 first2,
3   InputIterator3 stencil,         ForwardIterator result,
4   BinaryFunction binary_op,      Predicate pred);
```

## Example use:

```
1 #include <thrust/transform.h>
2 #include <thrust/functional.h>
3
4 int input1[6] = {-5, 0, 2, 3, 2, 4};
5 int input2[6] = { 3, 6, -2, 1, 2, 3};
6 int stencil[8] = { 1, 0, 1, 0, 1, 0};
7 int output[6];
8
9 thrust::plus<int> op;
10 thrust::identity<int> identity;
11
12 thrust::transform_if(input1, input1 + 6, input2, stencil,
13                     output, op, identity);
14
15 // output is now {-2, 0, 0, 3, 4, 4};
```



# Histogram Example

---

Given the sequence:

[2 1 0 0 2 2 1 1 1 1 4]

the dense histogram would be:

[2 5 3 0 1]

the sparse histogram would be:

[(0,2), (1,5), (2,3), (4,1)]

- How might we implement these operations with thrust algorithms?
- The first step - sort!
- Even when a serial algorithm might not involve sorting it is often a useful primitive when using thrust
- Let's examine the sparse case - we can do it with one thrust call!





# Segmented Reduction

```
1 thrust::pair<OutputIterator1,OutputIterator2>  
2 thrust::reduce_by_key(InputIterator1 keys_first,  
3                      InputIterator1 keys_last,  
4                      InputIterator2 values_first,  
5                      OutputIterator1 keys_output,  
6                      OutputIterator2 values_output);
```

We perform multiple reductions using adjacent identical keys to determine which reductions to perform

Keys : 1 3 3 3 2 2 1  
Vals : 9 8 7 6 5 4 3

Out Keys : 1 3 2 1  
Out Vals : 9 21 9 3

So how do we use this to do the histogram?



# Sparse Histogram

After sorting...

Sequence:           0 0 1 1 1 1 1 2 2 2 4 (keys)

Const. Iterator: 1 1 1 1 1 1 1 1 1 1 1 (vals)

Out Keys: 0 1 2 4 E X X X X X X

Out Vals: 2 5 3 1 E X X X X X X

```
1 thrust::device_vector<int> data(11); // = [2 1 0 0 2 2 1 1 1 1 4]
2 thrust::device_vector<int> histogram_values(11);
3 thrust::device_vector<int> histogram_counts(11);
4
5 thrust::sort(data.begin(), data.end());
6
7 typedef thrust::device_vector<int>::iterator devIt;
8 thrust::pair<devIt, devIt> endIterators =
9     thrust::reduce_by_key(data.begin(), data.end(),
10                          thrust::make_constant_iterator(1),
11                          histogram_values.begin(),
12                          histogram_counts.begin());
13
14 int num_values = endIterators.first - histogram_values.begin();
```



# Sparse Histogram

---

- What if we don't want to over-allocate space for `histogram_values` and `histogram_counts`?
- Must count how many distinct values in keys
- Can do this with a clever use of `inner_product`
- Inner product:  $(a_0 \otimes b_0) \oplus (a_1 \otimes b_1) \oplus \dots$
- Dot product is  $\otimes = \times$  and  $\oplus = +$

Sequence A:            0 0 1 1 1 1 1 2 2 2 4 (keys)

Sequence B:            0 0 1 1 1 1 1 2 2 2 4 (keys)

What should our operations be?



# Sparse Histogram

```
1 int num_bins = thrust::inner_product(data.begin(), data.end() - 1,  
2                                     data.begin() + 1,  
3                                     (int)1, thrust::plus<int>(),  
4                                     thrust::not_equal_to<int>());
```

Sequence A:            0 0 1 1 1 1 1 2 2 2 4 (keys)

Sequence B:            0 0 1 1 1 1 1 2 2 2 4 (keys)

A != B            :            0 1 0 0 0 0 1 0 0 1

num\_bins        :            1 + 0+1+0+0+0+0+1+0+0+1 = 4



## Second Example - Point Binning

---

- First generate a random collection of 2d points
- Then bin these points into a 2d grid of cells
- Finally extract which cells have only one point in them

•		• •	
• • •	•		
		•	• •
• • •	•		



# Random Number Generation

## With Thrust

---

- Can do RNG on both host and device
- We will focus on host generation
- We need both a generator AND a distribution
- A Generator produces random bits
  - `thrust::default_random_engine`
- The Distribution turns these bits into something useful - uniformly distributed floats between  $[-3, 10]$  for example
  - `thrust::uniform_real_distribution`
  - `thrust::uniform_int_distribution`



# Random Point Generation

---

```
1  // return a random vec2 in  $[0,1]^2$ 
2  vec2 make_random_float2(void)
3  {
4      //The static is important!
5      static thrust::default_random_engine rng;
6      static thrust::uniform_real_distribution<float> u01(0.0f, 1.0f);
7      float x = u01(rng);
8      float y = u01(rng);
9      return vec2(x,y);
10 }
11
12 thrust::host_vector<float2> h_points(N);
13 thrust::generate(h_points.begin(), h_points.end(), make_random_float2);
14
15 thrust::device_vector<float2> points = h_points;
```



# Grid Structure

---

```
1 // allocate storage for a 2D grid
2 // of dimensions w x h
3 unsigned int w = 200, h = 100;
4
5 // the grid data structure keeps a range per grid bucket:
6 // each bucket_begin[i] indexes the first element of
7 // bucket i's list of points
8 thrust::device_vector<unsigned int> bucket_begin(w*h);
9
10 // each bucket_end[i] indexes one past the last element of
11 // bucket i's list of points
12 thrust::device_vector<unsigned int> bucket_end(w*h);
13
14 // allocate storage for each point's bucket index
15 thrust::device_vector<unsigned int> bucket_indices(N);
```





# Point to Bucket Functor

```
1  // hash a point in the unit square to the index of
2  // the grid bucket that contains it
3  struct point_to_bucket_index :
4      thrust::unary_function<float2,unsigned int>
5  {
6      float width;  // buckets in the x dimension (grid spacing = 1/width)
7      float height; // buckets in the y dimension (grid spacing = 1/height)
8
9      __host__ __device__
10     point_to_bucket_index(unsigned int width, unsigned int height)
11         : width(width), height(height) {}
12
13     __host__ __device__
14     unsigned int operator()(const float2& v) const
15     {
16         // find the raster indices of p's bucket
17         unsigned int x = static_cast<unsigned int>(v.x * width);
18         unsigned int y = static_cast<unsigned int>(v.y * height);
19
20         // return the bucket's linear index
21         return y * width + x;
22     }
23
24 };
```



# Points → Cells

```
1  // transform the points to their bucket indices
2  thrust::transform(points.begin(),
3                    points.end(),
4                    bucket_indices.begin(),
5                    point_to_bucket_index(w,h));
6
7  // sort the points by their bucket index
8  thrust::sort_by_key(bucket_indices.begin(),
9                     bucket_indices.end(),
10                     points.begin());
```

Transform (assume w=10 h=10 here):

Points: (.7,.8) (.4,.1) (.6,.4) (.2,.3) (.62, .43)

Index :      87          14          64          32          64

Sort:

Index :      14          32          64          64          87

Points: (.4,.1) (.2,.3) (.6,.4) (.62, .43) (.7, .8)



# Determine Contents of Each Cell

---

- For each cell  $[0, w \times h)$  we need to figure out its bounds
- Hmm...makes me think of a `counting_iterator`
- The algorithms we need are `lower_bound` and `upper_bound`
- `lower_bound` takes a sequence and a list of search values
- For each search value, it finds the first place in the sequence it could be inserted without changing the ordering

Seq: 0 0 1 2 4 4 5 6 6 6 7

3 6

`lower_bound`: 3 6

Output: 4 7



# Determine Contents of Each Cell

---

- `upper_bound` is similar except it finds the *last* place in the sequence the value can be inserted without changing the ordering

Seq: 0 0 1 2 4 4 5 6 6 6 7  
                  3                  6

`upper_bound`: 3 6

Output: 4 10

Now we can determine the number of points in each cell by subtracting the output of `lower_bound` from `upper_bound`



# Code for Cell Determination

---

- The sequence we're searching *in* comes first
- The values we're searching *for* come next
- The output goes last

```
1 // find the beginning of each bucket's list of points
2 thrust::counting_iterator<unsigned int> search_begin(0);
3
4 thrust::lower_bound(bucket_indices.begin(),
5                     bucket_indices.end(),
6                     search_begin,
7                     search_begin + w*h,
8                     bucket_begin.begin());
9
10 // find the end of each bucket's list of points
11 thrust::upper_bound(bucket_indices.begin(),
12                    bucket_indices.end(),
13                    search_begin,
14                    search_begin + w*h,
15                    bucket_end.begin());
```



# Extracting Lonely Points

---

The function we need for this is called `remove_copy_if`

Cell:        0 1 2 3 4 5 6 7 8 9

# Points: 0 1 3 2 1 4 0 4 1 0

Out: 1 4 8

We need a predicate to determine which cells to remove based on the value of points

```
1  struct is_equal_to_one : thrust::unary_function<int, int>
2  {
3      __host__ __device__
4      int operator()(const int& v)
5      {
6          return v == 1;
7      }
8  };
```



# Extract Cells

---

```
1 thrust::device_vector<int> bucket_sizes(N);
2 thrust::transform(bucket_end.begin(), bucket_end.end(),
3                   bucket_begin.begin(), bucket_sizes.begin(),
4                   thrust::minus<int>());
5
6 int num_lonely_cells = thrust::count_if(bucket_sizes.begin(),
7                                         bucket_sizes.end(),
8                                         is_equal_to_one());
9
10 thrust::device_vector<int> lonely_cells(num_lonely_cells);
11 thrust::remove_copy_if(make_counting_iterator(0),
12                       make_counting_iterator(w*h),
13                       bucket_sizes.begin(),
14                       lonely_cells.begin(),
15                       is_equal_to_one() );
```



# Maximum Number of Points

---

To determine the cell with the most points, we could use `max_element` which returns an *iterator* to the largest element

```
1 thrust::device_vector<int>::iterator maxIt;  
2 maxIt = thrust::max_element(bucket_sizes.begin(), bucket_sizes.end());  
3  
4 int maxNum = *maxIt;  
5 int maxPos = maxIt - bucket_sizes.begin();
```

# Points: 0 1 3 2 1 4 0 4 1 0

maxNum: 4

maxPos: 5 (returns the first if there are duplicates)





# Vigenère Cipher

---

Plain text: ILIKEMYTEACHER

KEY: NOTNOTNOTNOTNO

=====

Cipher text: VZBXSFLHXNQARF

- By using multiple shifts (or permutations) instead of just 1 as in a substitution cipher the frequency distribution of the cipher text is obscured
- If we knew the key length, then we could solve multiple substitution ciphers

Plain text: ILIKEMYTEACHER

KEY: NOTNOTNOTNOTNO

=====

Cipher text: VZBXSFLHXNQARF



# Determine Key Length

First define an Index of Coincidence (IOC) between two texts as:

$$\frac{26 * \sum_{i=0}^N A_i == B_i}{N}$$

```
MYTEACHERISAWESOME
ILOVETHRUSTCODING
=====
000000100000000000
```

$$\text{IOC} = 26 * 1 / 17 = 1.53$$

- Defined so that the IOC between two *randomly* chosen texts is 1
- In English the IOC between two different texts (Moby Dick and The Great Gatsby) is  $\sim 1.73$
- Which is because letters are not chosen randomly, but have a non-uniform frequency distribution



# Breaking the Vigenère

---

Shifts don't line up and the IOC is  $\sim 1$ , the texts appear random

VZBXSFLHXNQARF  
VZBXSFLHXNQARF

VZBXSFLHXNQARF  
VZBXSFLHXNQARF

Shifts line up and suddenly the IOC jumps to  $\sim 1.73$ , because now at each position the "alphabet" though jumbled, is the same

VZBXSFLHXNQARF  
VZBXSFLHXNQARF

